

PROJECT-2: ANALYSIS REPORT OF DIABETES PATIENTS

Objective:

The main objective of this project is to diagnostically predict whether a patient has diabetes or not.

Note: The dataset used for this analysis is sourced from the National Institute of Diabetes and Digestive and Kidney Diseases.

About the dataset:

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases.

- 1. Pregnancies:** This variable represents the number of times the patient has been pregnant.
- 2. Glucose:** It is the plasma glucose concentration measured in milligrams per deciliter (mg/dL) of blood. This is a key indicator for diabetes diagnosis.
- 3. Blood Pressure:** This variable represents the diastolic blood pressure (mm Hg) of the patient.
- 4. Skin Thickness:** It indicates the skin thickness (mm) at the triceps area. Skin thickness can sometimes be relevant in diabetes diagnosis.
- 5. Insulin:** This variable represents the serum insulin level (μ U/ml). Insulin is a hormone that regulates blood sugar levels, and its measurement can be important in diabetes diagnosis.
- 6. BMI (Body Mass Index):** BMI is calculated from the weight and height of the patient. It's a measure of body fat and is often used to assess whether a person is underweight, normal weight, overweight, or obese.
- 7. Diabetes Pedigree Function:** This function is used to represent the diabetes pedigree function, which provides information about diabetes mellitus history in relatives and genetic influence.
- 8. Age:** This variable represents the age of the patient in years.
- 9. Outcome:** This is the target variable, and it indicates whether the patient has diabetes or not. It is binary, with values 0 and 1, where:
 - 0 typically indicates that the patient does not have diabetes.
 - 1 typically indicates that the patient has diabetes.

Steps:

Step-1: load the dataset and import necessary libraries

```
#importing libraries and loading csv file
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
data = pd.read_csv("diabetes.csv")
```

Step-2: Data Exploration

```
: #Get the summary statistics
print(data.describe())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Step-3: Data Pre-processing

```
#data preprocessing  
#check for the missing values  
print(data.isnull().sum())  
data = data.dropna()
```

```
Pregnancies      0  
Glucose          0  
BloodPressure    0  
SkinThickness    0  
Insulin          0  
BMI              0  
DiabetesPedigreeFunction  0  
Age              0  
Outcome          0  
dtype: int64
```

#Hence, there are no missing values in the data


```
#checking for duplicate data  
print(data[data.duplicated()])
```

```
Empty DataFrame  
Columns: [Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI, DiabetesPedigreeFunction,  
Age, Outcome]  
Index: []
```

#There are no duplicates in the data

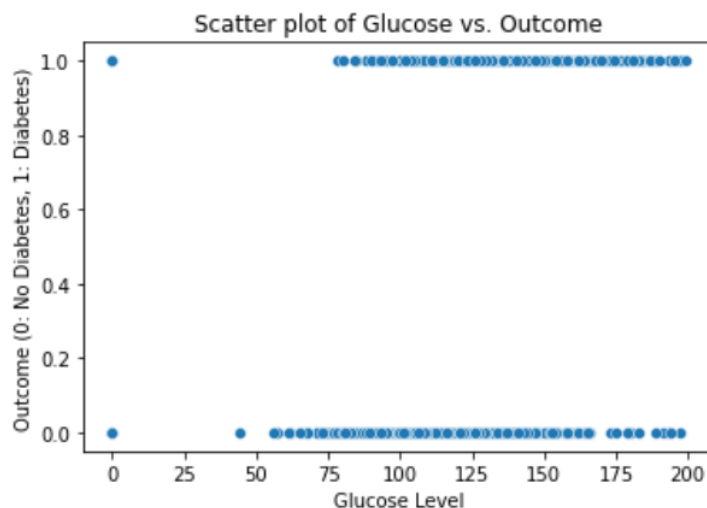
Step-4: Data Visualization

```
#Create histograms for numeric variables
sns.histplot(data['Glucose'], bins=20, kde=True)
plt.xlabel('Glucose Level')
plt.ylabel('Count')
plt.title('Distribution of Glucose Levels')
plt.show()
```

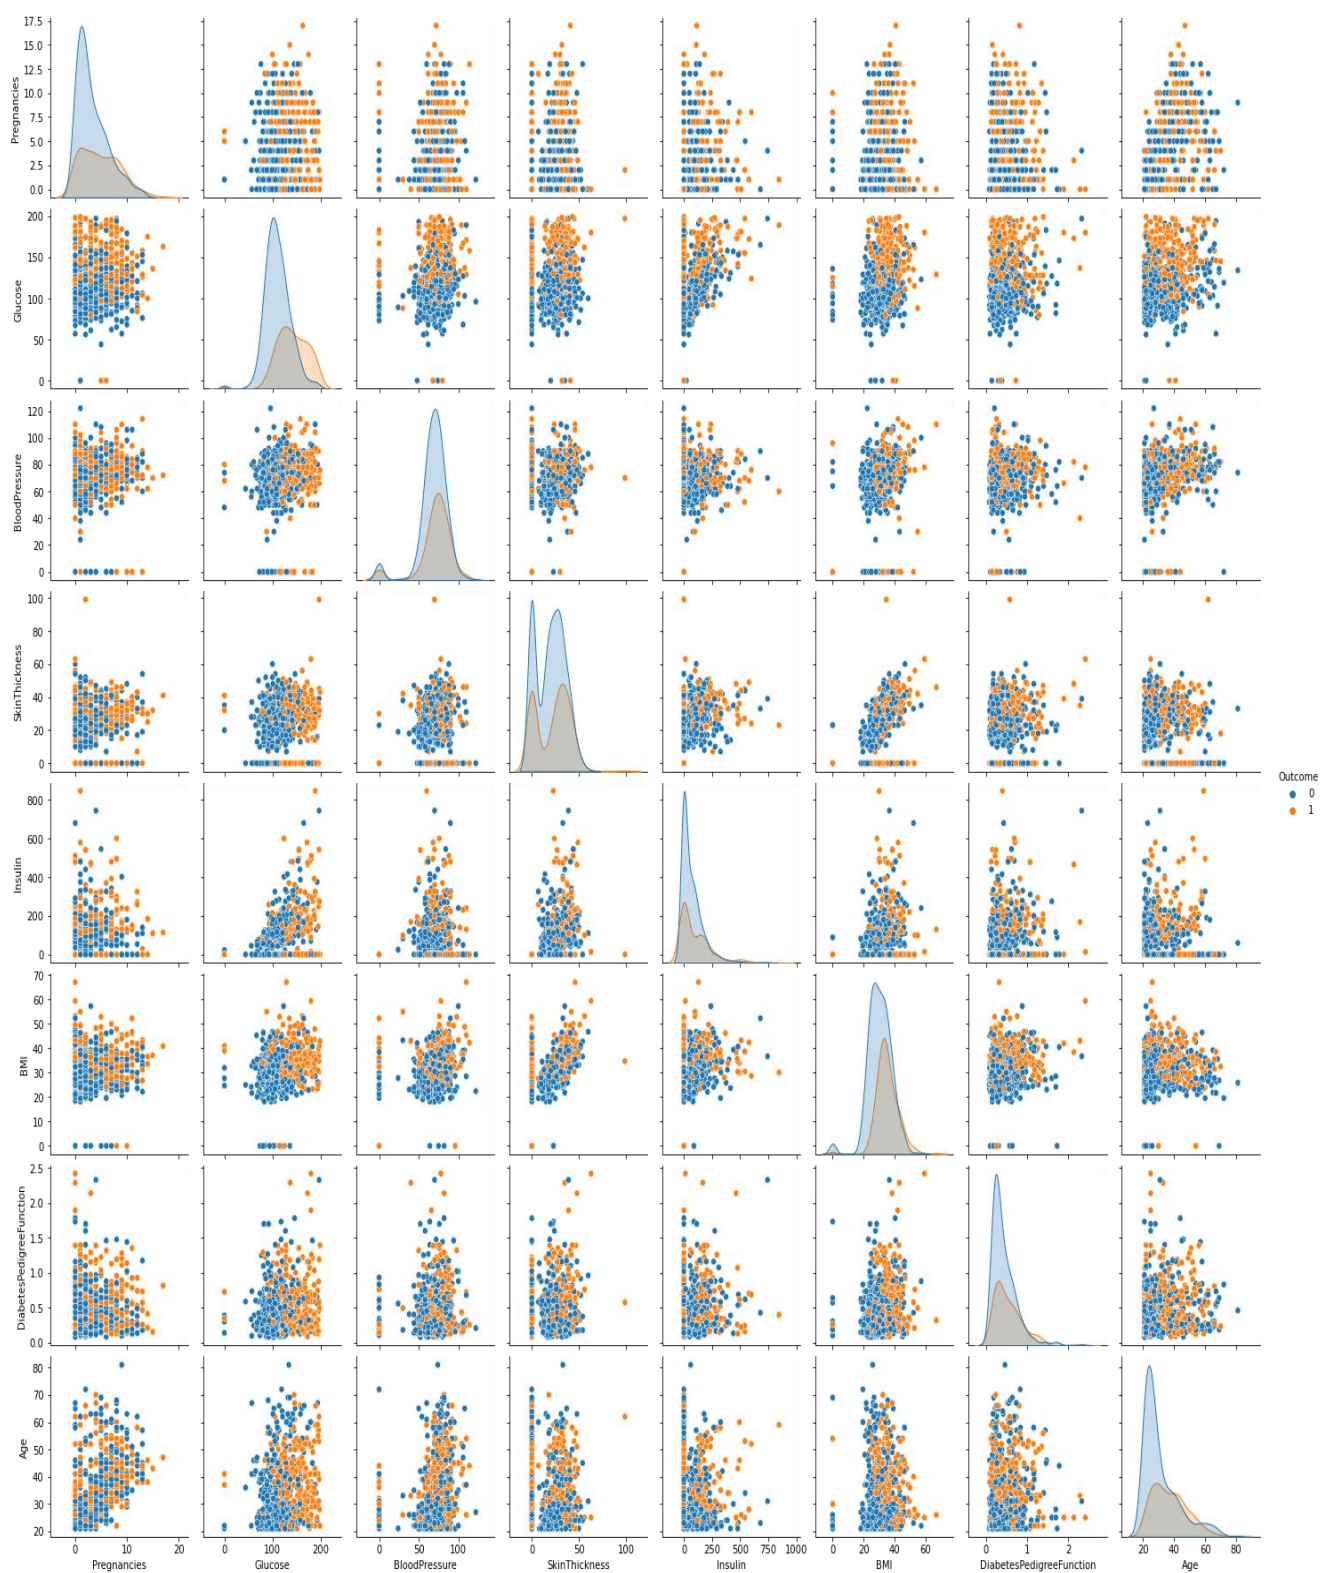


Step-5: Feature analysis


```
In [81]: # Scatter plot between Glucose and Outcome
sns.scatterplot(x='Glucose', y='Outcome', data=data)
plt.xlabel('Glucose Level')
plt.ylabel('Outcome (0: No Diabetes, 1: Diabetes)')
plt.title('Scatter plot of Glucose vs. Outcome')
plt.show()
```



```
: #creating visualization to understand the data better
sns.pairplot(data, hue="Outcome")
plt.show()
```



```
#Correlation matrix
correlation_matrix = data.corr()
sns.heatmap(correlation_matrix, annot=True)
plt.show()
```



Step-6: Model building and evaluation

```
#splitting the data into features(x) and target (y)
x = data.drop("Outcome", axis=1)
y = data["Outcome"]
```

```
#importing the libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
#splitting into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
#feature scaling
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

```
# training machine learning model using logistic regression
model = LogisticRegression()
model.fit(x_train, y_train)
```

```
LogisticRegression()
```



```
#evaluate the model-accuracy,precision,recall and F1 Score
```

```
#make predictions in the test set  
y_pred= model.predict(x_test)
```

```
#evaluate the model  
accuracy= accuracy_score(y_test, y_pred)
```

```
#classification report  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.80	0.81	99
1	0.65	0.67	0.66	55
accuracy			0.75	154
macro avg	0.73	0.74	0.73	154
weighted avg	0.76	0.75	0.75	154

```
#calculating the performance metrics  
accuracy = accuracy_score(y_test, y_pred)  
precision = precision_score(y_test, y_pred)  
recall= recall_score(y_test, y_pred)  
f1= f1_score(y_test, y_pred)
```

```
#report model performance  
print(f'Accuracy: {accuracy:.2f}')
```

```
print(f'Precision: {precision:.2f}')
```

```
print(f'Recall: {recall:.2f}')
```

```
print(f'F1-score: {f1:.2f}')
```

```
Accuracy: 0.75  
Precision: 0.65  
Recall: 0.67  
F1-score: 0.66
```

Conclusion:

- It has a decent level of precision, indicating that when it predicts positive cases (diabetes), it's correct about 65% of the time.
- The recall value indicates that the model is reasonably effective at identifying actual positive cases, capturing about 67% of them.
- The F1-score of 0.66 suggests that the model provides a balanced performance in terms of precision and recall.
- Out of a total of 768 patients, 268 have been diagnosed with diabetes.
- A higher number of pregnancies is associated with a decreased likelihood of diabetes.
- Patients with above-average blood pressure tend to have a lower likelihood of diabetes.
- An increase in blood pressure, BMI, and skin thickness is correlated with an increased likelihood of developing diabetes.
- Rising levels of glucose and insulin are linked to an increased risk of diabetes.